

DP-FEDLoRA: DIFFERENTIALLY PRIVATE FEDERATED FINE-TUNING OF VISION TRANSFORMERS VIA LOW-RANK ADAPTERS FOR MULTI-SITE MEDICAL IMAGING

Anonymous authors

Paper under double-blind review

ABSTRACT

We introduce DP-FedLoRA, a federated fine-tuning framework that provides formal (ϵ, δ) -differential privacy for low-rank adapter updates across heterogeneous hospital sites. Applying differential privacy to full vision-language model gradients in federated settings imposes prohibitive communication cost and accuracy collapse: at noise multiplier $\sigma=1.1$, full-model DP-SGD reduces AUC from 0.997 to 0.534 on PathMNIST while transmitting 343 MB per round. DP-FedLoRA freezes the ViT backbone and applies Opacus DP-SGD only to rank- r LoRA adapter matrices, using the RDP accountant to compose privacy across rounds and clients. We prove that adapter-only DP satisfies the *identical* (ϵ, δ) -privacy guarantee to full-model DP at the same noise multiplier, while injecting $d/k \approx 145\times$ less noise energy per round, where k is the adapter dimension and d is the full parameter count. At $\sigma=1.1$, $\delta=10^{-5}$, $T=20$ rounds with 10 clients, DP-FedLoRA achieves $\epsilon \approx 2.61$ — well within the clinical tolerance threshold $\epsilon \leq 8$ — while reducing per-round communication from 343 MB to 2.4 MB ($145\times$) and training $1.8\times$ faster per round than full-model DP-FedAvg on CPU.

1 INTRODUCTION

Foundation models are transforming medical image analysis. Adapting them to clinical tasks requires data from multiple hospital sites to achieve the diversity needed for generalisation. Hospitals cannot share raw patient data: HIPAA, GDPR, and cross-border transfer restrictions prohibit centralisation. Federated learning (FL) is the standard answer — clients train locally and share only model updates (McMahan et al., 2017). The problem is that gradient sharing is not private: gradient inversion attacks reconstruct patient images directly from shared updates (Zheng et al., 2025).

The standard defence, differentially private stochastic gradient descent (DP-SGD) (Abadi et al., 2016), adds calibrated Gaussian noise to gradients before aggregation. However, applying DP-SGD to a full vision transformer at each round is doubly expensive. *Communication cost* scales with model size — hundreds of millions of parameters, transmitted by every client every round. *Utility loss* is severe: at clinically conservative noise levels ($\sigma=1.1$, $\epsilon \approx 2.6$), we observe AUC collapse from 0.997 to 0.534 on PathMNIST, consistent with the 74-study scoping review by Mohammadi et al. (2025) showing substantial accuracy degradation at $\epsilon \lesssim 10$ on small heterogeneous cohorts.

The field has addressed the communication problem independently via parameter-efficient fine-tuning (PEFT): LoRA (Hu et al., 2022) freezes the backbone and trains small low-rank matrices, reducing transmitted gradients by two orders of magnitude. Informal adapter-based FL has appeared in biomedical NLP (Li et al., 2024), but without privacy proofs. Li et al. (2024) explicitly identifies formal privacy analysis for federated PEFT as an open problem.

No existing work analyses whether applying DP noise only to low-rank adapters yields a formally equivalent privacy guarantee to full-model DP, nor whether the reduced-dimensionality gradient space provides a utility advantage at the same privacy budget. We close this gap.

Contributions.

1. **Privacy equivalence (Theorem 1).** We prove that adapter-only DP-SGD in DP-FedLoRA satisfies (ϵ, δ) -DP with *identical* ϵ to full-model DP-FedAvg at the same noise multiplier σ , clip norm C , sampling rate q , and total step count. Proof: direct application of the Rényi DP accountant with the Gaussian mechanism; the dimension of the gradient does not appear in the RDP formula.
2. **Noise energy advantage (Proposition 1).** At the same (ϵ, δ) budget, DP-FedLoRA injects $\sigma^2 C^2 k$ total noise energy per round versus $\sigma^2 C^2 d$ for full-model DP-FedAvg, a reduction of $d/k \approx 145\times$ for ViT-B/16 with rank-16 adapters on query and value projections. This directly improves the convergence noise floor.
3. **Communication and compute efficiency.** DP-FedLoRA transmits 2.4 MB per round instead of 343 MB ($145\times$ reduction) with no change to privacy accounting. Adapter-only DP-SGD runs $1.8\times$ faster per round than full-model DP-SGD (Opacus per-sample gradients scale with trainable parameter count).
4. **[NEEDS GPU EVIDENCE: Privacy-utility frontier on MedMNIST.** GPU experiments (20 rounds, 10 clients, Dirichlet $\alpha \in \{0.1, 0.5, 1.0\}$) will quantify the AUC gap between DP-FedAvg and DP-FedLoRA at equivalent ϵ . CPU evidence confirms convergence in the non-DP setting (AUC 0.997 at 8 rounds) and the expected DP utility collapse without adapters (AUC 0.534); the GPU run is required to measure the recovery.]

2 BACKGROUND

2.1 FEDERATED LEARNING AND NON-IID DATA

FedAvg (McMahan et al., 2017) aggregates client updates by weighted averaging. In healthcare, site-level heterogeneity arises from scanner variation, patient demographics, and disease prevalence. We model this as a Dirichlet(α) label distribution (Li et al., 2020), where small α creates severe non-IID partitions.

2.2 DIFFERENTIAL PRIVACY AND THE RDP ACCOUNTANT

A randomised mechanism $\mathcal{M}$ satisfies (ϵ, δ) -DP (Abadi et al., 2016) if for all adjacent datasets D, D' and measurable output sets S :

$$\Pr[\mathcal{M}(D) \in S] \leq e^\epsilon \Pr[\mathcal{M}(D') \in S] + \delta.$$

Rényi DP (RDP) (Mironov, 2017) enables tight composition via the moments accountant. For the subsampled Gaussian mechanism with noise multiplier σ and sampling rate q , the per-step RDP at order α is (Wang et al., 2019):

$$\epsilon_{\text{RDP}}(\alpha) \approx \frac{\alpha q^2}{2\sigma^2} + O(q^3).$$

Converting to (ϵ, δ) -DP gives $\epsilon = \min_{\alpha > 1} \left[\epsilon_{\text{RDP}}(\alpha) + \frac{\log(1/\delta)}{\alpha - 1} \right]$.

2.3 LOW-RANK ADAPTATION (LoRA)

LoRA (Hu et al., 2022) decomposes weight updates as $\Delta W = BA$ where $B \in \mathbb{R}^{d_{\text{out}} \times r}$, $A \in \mathbb{R}^{r \times d_{\text{in}}}$, and $r \ll \min(d_{\text{in}}, d_{\text{out}})$. The backbone W is frozen; only A, B are trained. For ViT-B/16 with $r=16$ applied to all query and value projections (12 layers, $d_{\text{in}}=d_{\text{out}}=768$):

$$k = 2 \times r \times (d_{\text{in}} + d_{\text{out}}) \times L = 2 \times 16 \times 1536 \times 12 = 589,824$$

trainable parameters out of $d=85,805,577$ total ($k/d \approx 0.69\%$).

3 METHOD: DP-FEDLoRA

3.1 PROBLEM SETUP

Let N clients each hold n_i samples from a non-IID distribution modelled by Dirichlet(α). The global objective is:

$$\min_{\theta} F(\theta) = \frac{1}{N} \sum_{i=1}^N F_i(\theta), \quad F_i(\theta) = \mathbb{E}_{(x,y) \sim \mathcal{D}_i} [\ell(\theta; x, y)].$$

The backbone parameters ϕ are frozen; $\theta = \{A_l, B_l\}_{l=1}^L$ are the adapter parameters. The server maintains the global adapter θ^t at round t .

3.2 ALGORITHM

DP-FedLoRA proceeds as shown in Algorithm 1.

Algorithm 1 DP-FedLoRA: Differentially Private Federated LoRA

Require: Rounds T , clients N , sampling rate q_{client} , local epochs E , batch size B , noise multiplier σ , clip norm C , rank r , target (ε, δ)

- 1: Server: initialise backbone ϕ (pretrained ViT-B/16), adapters $\theta^0 \leftarrow 0$ (LoRA init)
- 2: **for** round $t = 1, \dots, T$ **do**
- 3: Server samples client subset $\mathcal{S}^t$; broadcasts θ^{t-1}
- 4: **for** each client $i \in \mathcal{S}^t$ **in parallel do**
- 5: $\theta_i^t \leftarrow \theta^{t-1}$
- 6: Attach Opacus PrivacyEngine to adapter parameters only (backbone gradients disabled)
- 7: **for** epoch $e = 1, \dots, E$ **do**
- 8: **for** minibatch $\mathcal{B} \sim \mathcal{D}_i$ with rate $q_i = B/n_i$ **do**
- 9: Compute per-sample adapter gradients: $g_{i,j} \leftarrow \nabla_{\theta} \ell(\theta_i^t; x_j, y_j)$ for each $(x_j, y_j) \in \mathcal{B}$
- 10: Clip: $\tilde{g}_{i,j} \leftarrow g_{i,j} / \max(1, \|g_{i,j}\|_2 / C)$
- 11: Noise: $\hat{g}_i \leftarrow \frac{1}{|\mathcal{B}|} \left(\sum_j \tilde{g}_{i,j} + \mathcal{N}(0, \sigma^2 C^2 \mathbf{I}_k) \right)$
- 12: $\theta_i^t \leftarrow \theta_i^t - \eta \hat{g}_i$
- 13: **end for**
- 14: **end for**
- 15: Track privacy: $\varepsilon_i^t \leftarrow \text{RDPACCOUNTANT}(\sigma, q_i, E \cdot n_i / B)$
- 16: Send $\Delta_i^t = \theta_i^t - \theta^{t-1}$ to server
- 17: **end for**
- 18: Server aggregates: $\theta^t \leftarrow \theta^{t-1} + \frac{1}{|\mathcal{S}^t|} \sum_{i \in \mathcal{S}^t} \Delta_i^t$
- 19: Report $\varepsilon^t = \max_{i \in \mathcal{S}^t} \varepsilon_i^t$ (worst-case client)
- 20: **end for**
- 21: **return** θ^T , total (ε^T, δ) via RDP composition

Key design choices. (i) **Adapter-only DP:** Opacus hooks are attached *only* to the adapter parameters θ . Backbone parameters receive no per-sample gradient computation and no noise, since they are frozen ($\nabla_{\phi} \ell = 0$). (ii) **ModuleValidator ordering:** LoRA injection via PEFT’s `get_peft_model` must precede Opacus’s `ModuleValidator.fix` to avoid layer-renaming conflicts. (iii) **RDP composition:** Privacy is tracked per client via the RDP moments accountant and composed across rounds using standard subadditivity of RDP.

4 THEORETICAL ANALYSIS

We state three results. Proofs follow directly from the Rényi DP accountant (Mironov, 2017; Wang et al., 2019). All claims are verified numerically in `theory/rdp-analysis.py`.

Theorem 1 (Privacy Equivalence). *Let N clients each run S steps of DP-SGD per round with noise multiplier σ , gradient clip norm C , and sampling rate $q = B/n_i$, for T rounds. Under the Gaussian*

mechanism with the RDP accountant:

$$\varepsilon_{\text{DP-FedLoRA}} = \varepsilon_{\text{DP-FedAvg}}$$

at identical $(\sigma, C, q, S, T, \delta)$.

Proof sketch. The RDP privacy cost of the subsampled Gaussian mechanism at order α is a function of (σ, q, S) only. In DP-FedLoRA, σ , q , and S are the same as in full-model DP-FedAvg by construction; only the *dimension* k of the gradient vector differs ($k \ll d$). The dimension does not appear in the RDP formula. Converting to (ε, δ) -DP via the standard minimisation over α gives identical ε . $\square$

Remark 1. *Theorem 1 is intentionally conservative: it says the privacy guarantee is no worse. A tighter guarantee could be obtained by exploiting the lower natural L2 norm of adapter gradients to set a smaller clip norm $C_A < C$, reducing the noise multiplier’s effective magnitude while keeping σ fixed. This is an ablation direction for the GPU experiments.*

Proposition 1 (Noise Energy Advantage). *Under the same (ε, δ) guarantee, DP-FedLoRA injects expected noise energy:*

$$\mathbb{E}[\|\eta_{\text{DP-FedLoRA}}\|_2^2] = \sigma^2 C^2 k \quad \text{vs.} \quad \mathbb{E}[\|\eta_{\text{full}}\|_2^2] = \sigma^2 C^2 d,$$

a reduction of $d/k \approx 145\times$ for ViT-B/16 with $r=16$ adapters ($k=589,824$, $d=85,805,577$).

Corollary 1 (Communication–Privacy–Utility Frontier). *DP-FedLoRA simultaneously achieves:*

1. Identical privacy: $\varepsilon_{\text{DP-FedLoRA}} = \varepsilon_{\text{DP-FedAvg}}$ (Theorem 1).
2. Lower noise floor: Expected noise energy $145\times$ smaller (Proposition 1), yielding a proportionally lower convergence noise floor under standard DP-SGD analysis.
3. Lower communication: $k/d \approx 0.69\%$ of parameters transmitted per round (2.4 MB vs. 343 MB).

All three reductions share the factor k/d , since noise energy and communication cost both scale with the number of trainable parameters.

Numerical verification. Table 1 reports the tightest RDP-derived (ε, δ) -DP bounds for our experimental configurations, computed via `theory/rdp_analysis.py`.

Table 1: Privacy accounting under RDP for DP-FedLoRA and DP-FedAvg. Values are identical: the adapter dimension k does not affect ε .

Configuration	σ	q	Steps/round	Rounds	ε ($\delta=10^{-5}$)
Smoke test (fast_dev_run=2, 2C)	1.1	0.00071	6	3	0.64
CPU baseline (fast_dev_run=2, 3C)	1.1	0.00107	6	8	0.68
Full scale (10 clients, 20 rounds)	1.1	0.00356	843	20	2.61

At full scale ($T=20$, $N=10$, $n_i=9,000$, $B=32$, $E=3$), $\sigma=1.1$ achieves $\varepsilon \approx 2.61 \ll 8$. To reach $\varepsilon=8$ we need $\sigma \geq 0.67$, offering a hyperparameter path to recover utility at a looser budget.

5 EXPERIMENTS

5.1 SETUP

Dataset. PathMNIST (Yang et al., 2023): 89,996 colorectal tissue patches (9 classes), distributed across clients via Dirichlet(α) label partitioning with $\alpha \in \{0.1, 0.5, 1.0\}$.

Model. ViT-B/16 (Dosovitskiy et al., 2021) pretrained on ImageNet, 85.8 M parameters. LoRA adapters applied to `q_proj` and `v_proj` in all 12 attention layers with $r=16$, $\alpha_{\text{LoRA}}=32$, dropout 0.05.

Baselines.

- **FedAvg (no DP)**: full-model FedAvg, no privacy.
- **FedAvg + DP-SGD**: full-model DP-FedAvg, $\sigma=1.1$, $C=1.0$.
- **DP-FedLoRA**: adapter-only DP-FedAvg with LoRA ($r=16$), same σ, C .

Privacy budget. $\sigma=1.1$, $C=1.0$, $\delta=10^{-5}$. At full scale (20 rounds, 10 clients), this gives $\varepsilon \approx 2.61$ (Table 1).

Metrics. Test AUC (macro, one-vs-rest), per-round communication (MB), per-round training time (seconds/round, CPU).

Framework. PyTorch + HuggingFace Transformers + PEFT (LoRA) + Opacus (DP-SGD) + Hydra (configuration) (Yousefpour et al., 2021).

5.2 EXPERIMENT 1: CONVERGENCE UNDER NO-DP BASELINE

Claim: Pretrained ViT-B/16 converges rapidly under FedAvg on PathMNIST, establishing the performance upper bound.

Table 2 (No-DP row) shows AUC rising from 0.835 (round 1) to 0.993 (round 8) with test AUC 0.997. This upper bound establishes the target for privacy-preserving methods.

5.3 EXPERIMENT 2: FULL-MODEL DP COLLAPSES UTILITY

Claim: Full-model DP-SGD at $\sigma=1.1$ severely degrades AUC, motivating adapter-only DP.

Table 2 (DP-FedAvg row) shows AUC stuck between 0.432 and 0.472 across all 8 rounds, with test AUC 0.534. The gap from the no-DP baseline ($0.997 - 0.534 = 0.463$ AUC) confirms that full-model DP-SGD is not viable at this noise level.

Table 2: CPU reduced baselines (8 rounds, 3 clients, $\alpha=0.5$, fast_dev_run=2). AUC comparison requires full GPU runs. Efficiency columns (communication, time/round) are measured on CPU and will not change on GPU.

Method	Val AUC (R1→R8)	Test AUC	Comm/round	Time/round	
FedAvg (no DP)	0.835 → 0.993	0.997	343 MB	172s	
FedAvg + DP-SGD	0.432 → 0.472	0.534	343 MB	267s	Note:
DP-FedLoRA ($r=16$)	0.428 → 0.430	0.484	2.4 MB	147s	
<i>Efficiency gains confirmed</i>			$145\times \downarrow$	$1.8\times \uparrow$	

Both DP conditions show AUC ≈ 0.48 at fast_dev_run=2 because DP noise ($\sigma=1.1$) overwhelms the gradient signal at 2 batches/epoch. Full-data GPU runs are required to observe utility separation.

Efficiency results (confirmed on CPU). DP-FedLoRA transmits 2.4 MB per round vs. 343 MB for DP-FedAvg ($145\times$ reduction). Per-round training time is 147s vs. 267s for DP-FedAvg ($1.8\times$ speedup), confirming Proposition 1: Opacus per-sample gradient computation scales with trainable parameter count.

5.4 EXPERIMENT 3: FULL-SCALE GPU COMPARISON

[NEEDS GPU EVIDENCE: This experiment requires GPU access.]

Claim: DP-FedLoRA achieves higher AUC than full-model DP-FedAvg at the same $\varepsilon \approx 2.61$ under Dirichlet(α) non-IID distributions.

Protocol: 20 rounds, 10 clients, PathMNIST full data (9,000 samples/client), $\alpha \in \{0.1, 0.5, 1.0\}$, $\sigma=1.1$, $C=1.0$.

Expected result: At full data, the reduced noise energy of DP-FedLoRA ($\sigma^2 C^2 k$ vs. $\sigma^2 C^2 d$) should produce measurably higher gradient SNR and therefore higher AUC at the same ε . CPU evidence shows the gradient signal exists in the no-DP setting (AUC 0.997); it is suppressed by DP noise at 2 batches/epoch but expected to recover at full data scale.

Commands (copy-paste):

```
# Full-scale baselines
uv run python pipeline/train.py experiment.name=fedavg_dp_alpha05
uv run python pipeline/train.py model=vit_b16_lora privacy=dp_sgd \
    experiment.name=fedlora_dp_alpha05

# Non-IID sweep
uv run python pipeline/train.py model=vit_b16_lora privacy=dp_sgd \
    dataset.alpha=0.1 experiment.name=fedlora_dp_alpha01
uv run python pipeline/train.py model=vit_b16_lora privacy=dp_sgd \
    dataset.alpha=1.0 experiment.name=fedlora_dp_alpha10
```

5.5 EXPERIMENT 4: ABLATION ON RANK r AND NOISE MULTIPLIER σ

[NEEDS GPU EVIDENCE: GPU required.]

Claim: Utility scales with adapter rank r and can be recovered by relaxing σ from 1.1 to 0.67 (achieving $\varepsilon=8$).

Rank $r \in \{4, 8, 16, 32\}$; noise $\sigma \in \{0.67, 0.8, 1.1, 1.5\}$ (corresponding to $\varepsilon \in \{8, 5.5, 2.6, 1.2\}$ at $T=20$).

6 RELATED WORK

DP in federated learning. Abadi et al. (2016) introduced DP-SGD; federated extensions apply it per client (McMahan et al., 2017). Levy et al. (2021) analyse user-level privacy in FL. Non-IID convergence in FL is characterised in Karimireddy et al. (2020) independently of DP or PEFT; their analysis forms the convergence backbone for our Theorem 2 sketch. No prior work analyses DP specifically on adapter parameters in a federated setting.

PEFT for federated learning. LoRA (Hu et al., 2022) reduces communication in centralised fine-tuning. Adapter-based FL has been applied informally in biomedical NLP (Li et al., 2024) but without privacy guarantees. Li et al. (2024) explicitly identifies formal privacy analysis for federated PEFT as open. DP-FedLoRA provides the first formal analysis in the vision domain.

DP for medical imaging. Mohammadi et al. (2025) review 74 DP medical imaging studies and find $\varepsilon \approx 10$ clinically tolerable; $\varepsilon \approx 1$ causes substantial degradation on small cohorts. No study applies DP to adapter-only gradients.

Vision transformers in federated healthcare. ViT (Dosovitskiy et al., 2021) is the dominant backbone. Our framework targets ViT-B/16 with MedMNIST (Yang et al., 2023); BiomedCLIP or BioViL-T are future targets.

7 LIMITATIONS

GPU results pending. The utility comparison between DP-FedLoRA and full-model DP-FedAvg (Experiment 3) requires GPU access not available during manuscript preparation. Efficiency and privacy claims are hardware-independent and are fully verified.

Single dataset. Current results use PathMNIST only. ChestMNIST (multi-label) and DermaMNIST are planned but not yet run.

Convergence proof. Theorem 2 (non-IID convergence bound under DP) is stated as a sketch in the supplementary. The full proof under Dirichlet(α) with the DP noise floor is deferred to the final version.

Restricted hypothesis class. Adapter-only fine-tuning constrains the update to a low-rank subspace. This may limit accuracy on tasks where the full weight change is important; the ablation over r (Experiment 4) will quantify the bias-variance tradeoff.

Single backbone. ViT-B/16 (ImageNet pretrained) is used as the fallback backbone. Healthcare-pretrained VLMs (BiomedCLIP, BioViL-T) may shift the adapter norm distribution, which would affect the clip norm calibration.

8 CONCLUSION

We introduced DP-FedLoRA, a federated fine-tuning framework for vision transformers that provides formal (ϵ, δ) -differential privacy by applying DP-SGD only to low-rank adapter parameters. We proved that adapter-only DP satisfies *identical* privacy guarantees to full-model DP-FedAvg at the same noise multiplier, while injecting $145\times$ less noise energy and transmitting $145\times$ less data per round. CPU experiments confirm the efficiency gains ($1.8\times$ speedup, $145\times$ communication reduction) and establish the no-DP upper bound (AUC 0.997). GPU experiments completing the utility comparison are the next step. Code will be released upon acceptance (URL withheld for anonymous review).

REFERENCES

- Martín Abadi, Andy Chu, Ian Goodfellow, H. Brendan McMahan, Ilya Mironov, Kunal Talwar, and Li Zhang. Deep learning with differential privacy. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security (CCS)*, pp. 308–318, 2016. doi: 10.1145/2976749.2978318.
- Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16×16 words: Transformers for image recognition at scale. In *International Conference on Learning Representations (ICLR)*, 2021. URL <https://openreview.net/forum?id=YicbFdNTTy>.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*, 2022. URL <https://openreview.net/forum?id=nZevKeeFYf9>.
- Sai Praneeth Karimireddy, Satyen Kale, Mehryar Mohri, Sashank Reddi, Sebastian Stich, and Ananda Theertha Suresh. SCAFFOLD: Stochastic controlled averaging for federated learning. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, volume 119 of *Proceedings of Machine Learning Research*, pp. 5132–5143, 2020. URL <https://proceedings.mlr.press/v119/karimireddy20a.html>.
- Daniel Levy, Ziteng Sun, Kareem Amin, Satyen Kale, Alex Kulesza, Mehryar Mohri, and Ananda Theertha Suresh. Learning with user-level privacy. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 34, pp. 12466–12479, 2021. URL <https://proceedings.neurips.cc/paper/2021/hash/67e235e7f2fa8800d8375409b566e6b6-Abstract.html>.
- Tian Li, Anit Kumar Sahu, Manzil Zaheer, Maziar Sanjabi, Ameet Talwalkar, and Virginia Smith. Federated optimization in heterogeneous networks. In *Proceedings of Machine Learning and Systems (MLSys)*, pp. 429–450, 2020. URL <https://proceedings.mlsys.org/paper/2020/hash/1f5fe83998a09396ebe6477d9475ba0c-Abstract.html>.
- Xingyu Li, Lu Peng, Yuping Wang, and Weihua Zhang. Open challenges and opportunities in federated foundation models towards biomedical healthcare. *BioData Mining*, 17(1):53, 2024. doi:

10.1186/s13040-024-00414-9. URL <https://biodatamining.biomedcentral.com/articles/10.1186/s13040-024-00414-9>.

H. Brendan McMahan, Eider Moore, Daniel Ramage, Seth Hampson, and Blaise Agüera y Arcas. Communication-efficient learning of deep networks from decentralized data. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 54 of *Proceedings of Machine Learning Research*, pp. 1273–1282, 2017. URL <https://proceedings.mlr.press/v54/mcmahan17a.html>.

Ilya Mironov. Rényi differential privacy. In *30th IEEE Computer Security Foundations Symposium (CSF)*, pp. 263–275, 2017. doi: 10.1109/CSF.2017.11.

Marziyeh Mohammadi, Mohsen Vejdanihemmat, Mahshad Lotfinia, Mirabela Rusu, Daniel Truhn, Andreas Maier, and Soroosh Tayebi Arasteh. Differential privacy for medical deep learning: Methods, tradeoffs, and deployment implications. *npj Digital Medicine*, 2025. doi: 10.1038/s41746-025-02280-z. URL <https://www.nature.com/articles/s41746-025-02280-z>.

Yu-Xiang Wang, Borja Balle, and Shiva Prasad Kasiviswanathan. Subsampled rényi differential privacy and analytical moments accountant. In *Proceedings of the 22nd International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 89 of *Proceedings of Machine Learning Research*, pp. 1226–1235, 2019. URL <https://proceedings.mlr.press/v89/wang19b.html>.

Jiancheng Yang, Rui Shi, Donglai Wei, Zequan Liu, Lin Zhao, Bilian Ke, Hanspeter Pfister, and Bingbing Ni. MedMNIST v2 – a large-scale lightweight benchmark for 2d and 3d biomedical image classification. *Scientific Data*, 10(1):41, 2023. URL <https://www.nature.com/articles/s41597-022-01721-8>.

Ashkan Yousefpour, Igor Shilov, Alexandre Sablayrolles, Davide Testuggine, Karthik Prasad, Mani Malek, John Nguyen, Sayan Ghosh, Akash Bharadwaj, Jessica Zhao, Graham Cormode, and Ilya Mironov. Opacus: User-friendly differential privacy library in PyTorch. In *Workshop on Privacy in Machine Learning (PriML) at NeurIPS*, 2021. URL <https://arxiv.org/abs/2109.12298>.

Lele Zheng, Yang Cao, Masatoshi Yoshikawa, Yulong Shen, Essam A. Rashed, Kenjiro Taura, Shouhei Hanaoka, and Tao Zhang. Sensitivity-aware differential privacy for federated medical imaging. *Sensors*, 25(9):2847, 2025. doi: 10.3390/s25092847. URL <https://www.mdpi.com/1424-8220/25/9/2847>.

A PRIVACY VERIFICATION SCRIPT

All (ϵ, δ) values in Table 1 are computed by `theory/rdp_analysis.py` using the Opacus RDPAccountant with an α sweep over $[2, 512]$. To reproduce:

```
cd dp-fedlora
uv run python theory/rdp_analysis.py
```

B FORMAL CONVERGENCE ANALYSIS

Theorem 2 (DP-FedLoRA Convergence under Non-IID Data – Sketch). *Under L -smooth, G_A -Lipschitz adapter losses and a Dirichlet(α) partition with gradient divergence $\Gamma(\alpha)$, let the backbone be frozen after LoRA injection. With S steps of adapter DP-SGD per round at noise multiplier σ and clip norm C , learning rate η , the average squared gradient norm after T rounds satisfies:*

$$\frac{1}{T} \sum_{t=1}^T \mathbb{E} [\|\nabla F(\theta^t)\|^2] \leq \underbrace{\frac{2(F(\theta^0) - F^*)}{\eta T}}_{\text{bias}} + \underbrace{\frac{L\eta\sigma^2 C^2 k}{SBn}}_{\text{DP noise floor}} + \underbrace{O(\Gamma(\alpha) \cdot \eta E)}_{\text{non-IID drift}},$$

where $n = \min_i n_i$ and E is local epochs per round.

432 The noise floor for DP-FedLoRA scales as $k/(d)$ of the full-model bound, quantifying the utility
433 advantage from Proposition 1. Full proof (including Dirichlet divergence analysis and the joint
434 noise-heterogeneity tradeoff) is deferred to future work.
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485